

BloomSpan: Memory-Efficient Maximal Frequent Substring Mining for Large-Scale Text

Rauf Aliev¹[0009-0006-2564-7742] and Joeran Beel¹[0000-0002-4537-5573]

University of Siegen, Siegen, Germany
r.aliev@gmail.com, joeran.beel@uni-siegen.de

Abstract. BloomSpan is a memory-efficient algorithm for mining Maximal Contiguous Frequent Phrases (MCFPs) from large text corpora. MCFPs are strictly contiguous token sequences satisfying a minimum document-support threshold; unlike general Sequential Pattern Mining, no gaps are permitted. BloomSpan combines a Counting Bloom Filter for probabilistic frequency estimation with prioritized seed generation and depth-first expansion directly on the corpus, avoiding recursive projected databases. Seed generation is provably free of false negatives; output equivalence with exact baselines is confirmed on all Gutenberg-8k subsets where any baseline completes. On the **Gutenberg-BookCorpus-Cleaned** corpus (up to 3.07 B tokens), BloomSpan achieves 11-15x faster wall-clock execution than FHK with up to 2x lower peak memory.

Keywords: Frequent substring mining · Maximal contiguous phrases · Suffix array · Counting Bloom Filter · Large-scale text corpora

1 Introduction

The identification of frequent patterns within large sequences is a cornerstone problem in data mining. General Sequential Pattern Mining (SPM) algorithms such as GSP [2] and PrefixSpan [21] allow arbitrary gaps between elements and treat each element as a set (itemset). While mathematically powerful, this generality is a liability in domains such as document template extraction, plagiarism detection, and phrase mining, where patterns are single tokens forming rigid, unbroken structures. Introducing gaps into such searches leads to a *combinatorial explosion* of meaningless candidates: given a sequence $\langle A, B, C \rangle$, an SPM miner considers the pattern present even when elements appear as $A \dots \text{gap} \dots B \dots \text{gap} \dots C$, generating orders of magnitude more candidates than the contiguous case requires. The Maximal Contiguous Frequent Phrase (MCFP) mining problem [7] captures the constrained setting where no gap is allowed and each element is a single token.

Practitioners navigating this problem face a dichotomy of extremes. *Stringology-native suffix-array miners* build on decades of index engineering starting with Manber and Myers' suffix array [17], Kasai et al.'s linear-time LCP construction [15], and the enhanced suffix array framework of Abouelhoda, Kurtz, and

Ohlebusch [1]; practical deployments further rely on Nong–Zhang–Chan’s SA-IS suffix array construction [19] and, for compressed settings, on the FM-index and related self-indexes surveyed by Navarro and Mäkinen [18]. Specifically for substring mining, the Fischer–Heun–Kramer algorithm (FHK) [8, 9], the Deferred Frequency Index (DFI [23]), and related work [16, 4, 10] solve MCFP in asymptotically optimal time but carry constant factors that become significant on large natural-language alphabets and permissive operating points. Recent work at SPIRE 2024 continues to refine this line, including online [13] and faster offline [20] computation of closely related net-frequency queries. BloomSpan complements this line by moving orthogonally: it targets a document-frequency constraint on a bounded-memory probabilistic substrate rather than an exact index, trading worst-case guarantees for empirical memory savings on large natural-language corpora. *Gap-aware sequence miners* (BIDE+ [22], VMSP [12]) solve a strictly more general problem and manage a much larger search space, incurring severe memory pressure when restricted to contiguous extensions.

We present *BloomSpan*, which occupies a distinct operating point between these families: a Counting Bloom Filter bypasses explicit candidate generation, while a depth-first expansion with a global occupancy bitmask bounds memory to $O(W + M)$ in the input size. Our main contributions are: (1) an algorithm whose runtime decomposes into an input-linear pre-processing component and a data-dependent expansion component, scaling near-linearly per token on a 3 billion-token natural-language corpus at strict thresholds ($\sigma \geq 20$); (2) a probabilistic seed-filtering mechanism that provably produces no false negatives.

2 Problem Formalization

Let Σ be a finite token alphabet. A document $d_i = (t_1, \dots, t_n)$ is an ordered token sequence. Given corpus $\mathcal{C} = \{d_1, \dots, d_m\}$ and support threshold σ , the task is to extract all MCFPs.

- **Phrase (P):** A **strictly contiguous** subsequence $(t_i, t_{i+1}, \dots, t_{i+k-1})$.
- **Support (df):** $df(P) = |\{d \in \mathcal{C} : P \text{ appears in } d\}|$.
- **Seed (S):** A frequent n -gram ($df(S) \geq \sigma$) serving as an expansion anchor.
- **MCFP:** A phrase P with $df(P) \geq \sigma$ such that no strict super-phrase P' containing P satisfies $df(P') \geq \sigma$.

3 Proposed Algorithm: BloomSpan

BloomSpan transitions from pattern-growth projection trees to a prioritized Depth-First Search (DFS) combined with probabilistic filtering. It operates in four phases.

3.1 Phase 1: Database Loading and Probabilistic Frequency Estimation

The corpus is loaded and exact 1-gram document frequencies $df(t)$ are computed. A Counting Bloom Filter (CBF) [5] (up to 512 MB) is populated with all n -grams

Algorithm 1 BloomSpan (top-level)**Require:** Corpus $\mathcal{C}$, support threshold σ , seed length n , minimum phrase length $\ell_{\min}$ **Ensure:** Set $\mathcal{M}$ of MCFPs

- 1: $df_1 \leftarrow$ exact 1-gram document frequencies over $\mathcal{C}$
- 2: $\hat{c} \leftarrow$ CBF populated with all length- n windows of $\mathcal{C}$
- 3: $\mathcal{S} \leftarrow \{S : \hat{c}(S) \geq \sigma \wedge \min_{t \in S} df_1(t) \geq \sigma\}$ $\triangleright$ seed candidates
- 4: $\mathcal{S}' \leftarrow$ merge $\mathcal{S}$ by exact df , keep those with $df \geq \sigma$
- 5: sort $\mathcal{S}'$ by $\Psi(S) = |S| \cdot df(S)$ desc., then length desc., then lex.
- 6: $occupied \leftarrow$ all-zeros bitmask over $\sum_d |d|$ positions
- 7: $\mathcal{M} \leftarrow \emptyset$
- 8: **for** each candidate $S \in \mathcal{S}'$ in order **do**
- 9: **if** all occurrences of S lie in $occupied$ positions **then continue**
- 10: **end if**
- 11: DFS-Expand($S, occupied, \mathcal{M}$) $\triangleright$ Algorithm 2
- 12: **end for**
- 13: **return** $\mathcal{M}$




Fig. 1: BloomSpan pipeline overview.

($n \geq 2$) via a linear pass, providing a space-efficient frequency sketch akin to the Count-Min Sketch [3].

3.2 Phase 2: Seed Candidate Generation and Prioritization

A parallel corpus scan identifies n -gram seeds. A candidate is pruned if (a) its CBF estimate is below σ , or (b) any constituent 1-gram has $df < \sigma$. The CBF maintains an estimate $\hat{c}(S)$ of the total corpus occurrence count $c(S) = \sum_{d \in \mathcal{C}} \#\{i : d[i..i+n-1] = S\}$. Hash collisions can only over-estimate this count, so $\hat{c}(S) \geq c(S)$. Since each document containing S contributes at least one occurrence to $c(S)$, we have $c(S) \geq df(S)$ trivially. Composing both inequalities, $\hat{c}(S) \geq df(S)$, so any S with true $df(S) \geq \sigma$ passes the filter $\hat{c}(S) \geq \sigma$. By the

anti-monotonicity of document frequency, every MCFP of length $\ell \geq n$ contains at least one length- n contiguous sub-sequence S with $df(S) \geq \sigma$; hence every truly frequent longer phrase remains reachable from a surviving seed. Together these two properties guarantee no truly frequent pattern is pruned at this stage. (We track $c(\cdot)$ rather than $df(\cdot)$ to avoid the per-document reset cost on a CBF; the inequality $c \geq df$ is sufficient for the guarantee, and exact df values are recomputed in $O(\text{occ})$ time during seed merging in Phase 2.). The CBF filter is loose with respect to df — a bigram occurring 20 times within a single document passes — so absolute pruning ratios depend on the corpus topology. On Realistic Dataset 2 at $\sigma = 20$, $n = 2$ (Section 4.7), the combined CBF-and-1-gram filter eliminates 98% of candidate n -grams; we report per-corpus pruning ratios in Section 4.7. Valid seeds are assigned a coverage score $\Psi(P) = |P| \times df(P)$ and globally sorted by it, with ties broken first by descending length and then lexicographically, so that larger, more frequent “backbone” patterns are processed first.

3.3 Phase 3: DFS Expansion and Early Pruning

Expansion is performed directly on the original corpus using a stack, avoiding projected databases. Before right-extension, an *Early Pruning* step checks whether all occurrences of a candidate share the same preceding token; if so, the candidate is inherently non-maximal and is immediately pruned.

3.4 Phase 4: Maximality Check and Redundancy Control

An *Inverted Index* keyed on each pattern’s rarest token reduces the maximality check from an $O(P \times l_p)$ scan to an $O(p_{\text{rare}} \times l_p)$ lookup ($p_{\text{rare}} \ll P$). If a candidate is confirmed maximal, BloomSpan also absorbs and removes any previously discovered sub-phrases now subsumed by the longer pattern. A global *occupancy bitmask* flags all token positions belonging to confirmed maximal patterns; any subsequent candidate whose occurrences lie entirely within already-flagged positions is skipped before expansion.

3.5 Complexity

Let W = total tokens, M = CBF size, C = valid seed occurrences, P = maximal patterns found.

Space: $O(W + M + C + P)$. The corpus and bitmask each require $O(W)$; the CBF uses exactly M bytes; seed and index structures require $O(C + P \cdot l_p)$.

Time: Let W = total tokens, n = seed length (a constant; $n = 2$ throughout), C = surviving seed occurrences after Phase 2, and P = MCFPs in the output. Phase 1 (CBF population and 1-gram counting) takes $\Theta(W)$. Phase 2 (seed gathering, parallel scan and sort) takes $\Theta(W + C \log C)$. Phase 3 (DFS expansion) is data-dependent: each expansion step over a candidate with s occurrences takes $O(s)$ for the early-pruning check and the next-token grouping,

and the inverted-index maximality check costs $O(p_{\text{rare}} \cdot l_p \cdot l_{\text{super}})$ per candidate, where p_{rare} is the number of MCFPs containing the candidate’s rarest token. We do not establish a worst-case bound on Phase 3; instead, Section 4.7 reports approximately constant per-token mining cost (0.021–0.027 $\mu\text{s}/\text{token}$, log-log slope 0.98) across a $28\times$ range of input sizes at $\sigma = 20$, with a transition to super-linear scaling once the MCFP set itself grows large.

4 Experimental Evaluation

4.1 Setup

All algorithms are evaluated as JVM-based implementations on the same machine (Ultra 9, 64 GB RAM, OpenJDK 21, -Xmx40g); evaluating BloomSpan, FHK, and DFI within a single language stack ensures that reported speedups reflect algorithmic differences rather than runtime artifacts. SPMF baselines (GST/SA+LCP, VMSP, BIDE+) use the SPMF framework [11] (v2.65); FHK and DFI are in-house Java implementations. Reported times are wall-clock from algorithm invocation to termination, measured identically by a common benchmark runner; includes corpus parsing and, for FHK/DFI, suffix and LCP array construction, excludes one-time conversion to the SPMF int-encoded format (before the timed region, identical across algorithms). Each (algorithm, N) configuration was executed for 6 runs; the 1st is discarded as JVM warm-up, and medians with inter-quartile ranges (IQR) over the remaining 5 are reported. We use medians because occasional GC pauses introduce positive skew; across all BloomSpan runs the relative time IQR stayed below 5% of the median. Source code: <https://github.com/raliev/bloomspan>.

4.2 Baselines

1. **FHK** — in-house Java implementation of Fischer–Heun–Kramer [7, 9]: SA-IS suffix array [19], Kasai LCP array [15], stack-based bottom-up LCP-interval traversal.
2. **DFI** — in-house Java implementation of the Deferred Frequency Index [23]: same SA-IS + Kasai preprocessing, top-down DFS with early subtree pruning and an $O(n \log n)$ RMQ sparse table.
3. **GST/SA+LCP** — suffix-array baseline from SPMF.
4. **VMSP** — maximal sequential pattern miner from SPMF [12]; included as a sanity check.
5. **BIDE+ (Contiguous Maximal)** — maximal sequence miner [22] restricted to contiguous extensions.

4.3 Datasets

We evaluate on two synthetic and two realistic datasets covering a wide range of corpus sizes and pattern topologies.

Table 1: Characterisation of Gutenberg-8k subsets. TTR = type/token ratio.

N	Size	Total tokens	Vocabulary	Tokens / document			TTR
				median	mean	max	
100	21.2 MB	3,913,260	263,784	18,376	39,132	162,961	0.0674
250	72.0 MB	13,272,683	591,817	38,618	53,090	259,595	0.0446
500	156.1 MB	28,438,195	1,022,479	38,445	56,876	467,524	0.0360
1,000	324.1 MB	58,853,177	1,900,504	42,945	58,853	763,654	0.0323
3,000	1,015.1 MB	183,409,628	4,616,291	46,030	61,136	763,654	0.0252
5,000	1.7 GB	305,326,775	6,536,695	45,412	61,065	832,335	0.0214
8,544	2.9 GB	535,259,155	9,113,577	44,733	62,647	1,662,371	0.0170

- **Synthetic Dataset 1 (Sparse Injections).** N documents of 5000 unique synthetic words each, with three short fixed sentences injected into a small random subset of documents ($\sigma = 3$, min. phrase length 3). Sizes: $N = 100$ – $30,000$.
- **Synthetic Dataset 2 (Dense Nested Overlaps).** 300 heavily overlapping sentences of incrementally decreasing length injected into a unique-word background ($N = 15,000$). Designed as a worst-case topology for BloomSpan’s redundancy control.
- **Realistic Dataset 1 (Gutenberg-8k).** First N files of `Fhrozen/gutenberg8k` [6] (8544 books, 535 M tokens). Subsets characterised in Table 1. Parameters: $\sigma = 10$, min. phrase length 5.
- **Realistic Dataset 2 (Gutenberg-BookCorpus-Cleaned).** First N files of the 66,871-book Gutenberg-BookCorpus-Cleaned release [14] (N up to 50,000, ≈ 3.07 B tokens). Parameters: $\sigma = 20$, min. phrase length 10. SPMF-based baselines exhaust the heap at $N = 1,000$ and are excluded.

4.4 Results: Synthetic Dataset 1

Figure 2 and Table 2 show results for Synthetic Dataset 1. All stringology-native miners (FHK, DFI, BloomSpan) scale approximately linearly in time and memory up to $N = 10,000$. However, a performance crossover occurs at $N = 12,000$ (0.6×10^8 tokens). Beyond this point, FHK becomes the most efficient algorithm in both dimensions. At $N = 30,000$, FHK completes in 15.3 s using 38.9 GB, whereas BloomSpan requires 38.2 s and reaches the heap limit at 40.3 GB. This shifts the previous observation: while BloomSpan is highly competitive on smaller synthetic sets, FHK’s LCP-interval traversal proves more robust for the flat frequency distribution of unique synthetic background words at large scales. However, this changes for the realistic dataset.




Fig. 2: Execution time (L) and peak memory (R) on Synthetic Dataset 1..

Table 2: Performance on Synthetic Dataset 1 (5000 w/d). Time in ms; Mem in MB.

N (Docs)	BloomSpan		FHK		DFI		GST/SA+LCP	
	Time	Mem	Time	Mem	Time	Mem	Time	Mem
1,000	1,047	1,893	625	1,908	713	2,773	931	1,903
10,000	7,852	17,339	6,151	20,920	7,704	20,525	10,041	15,548
12,000	10,115	20,405	7,299	21,249	8,782	21,716	13,091	14,953
20,000	15,772	31,351	10,446	25,332	15,384	32,711	18,727	35,356
30,000	38,240	40,360	15,373	38,949	OOM	OOM	75,348	35,603

Table 3: Performance on Synthetic Dataset 2 at $N = 15,000$ (84M tokens). Medians with IQR over six runs. Time in ms; memory in MB.

Algorithm	Time (ms)	Memory (MB)
FHK	10,461 (401)	23,166 (587)
BloomSpan	12,589 (2,337)	26,809 (2,942)
DFI	14,861 (280)	32,903 (971)
BIDE+ (Contig. Maximal)	16,765 (318)	36,965 (866)


Fig. 3: Execution time (left) and peak memory (right) on Synthetic Dataset 2, $N = 15,000$, six runs. BIDE+ (non-maximal) and GST/SA+LCP exhaust the heap.

Table 4: Performance on Gutenberg-8k ($\sigma = 10$, min. phrase length 5). Time in ms; memory in MB. Medians over 10 runs.

N	BloomSpan		FHK		DFI		BIDE+ (Contig.)		GST/SA+LCP	
	Time	Mem	Time	Mem	Time	Mem	Time	Mem	Time	Mem
100	713	727	625	220	655	416	1,275	1,136	2,373	798
250	1,682	1,291	2,023	1,278	2,012	1,299	4,847	4,696	10,445	4,783
500	3,201	1,960	3,896	1,145	4,016	1,405	13,241	25,688	27,170	21,482
1,000	6,608	4,483	8,425	1,875	8,313	2,158	OOM	OOM	OOM	OOM
1,500	–	–	12,714	2,602	12,752	2,517	OOM	OOM	OOM	OOM
3,000	28,550	7,118	25,427	3,324	25,185	3,745	OOM	OOM	OOM	OOM
5,000	141,023	5,416	117,228	5,155	111,864	5,179	OOM	OOM	OOM	OOM
8,000	OOM	OOM	2,214,653	9,828	1,461,332	9,819	OOM	OOM	OOM	OOM

4.5 Results: Synthetic Dataset 2 (Dense Nested Overlaps)

Table 3 and Figure 3 show results at $N = 15,000$ for the adversarial dataset. FHK is fastest and most memory-efficient: its LCP-interval traversal directly exploits the dense overlapping suffix structure. BloomSpan is second-fastest; its occupancy-bitmask pruning outperforms BIDE+’s per-pattern redundancy tracking despite CBF filtering losing some of its advantage when nearly every n -gram is frequent. BIDE+ (Contiguous, non-maximal) and GST/SA+LCP both exhaust the 40 GB heap.

4.6 Results: Realistic Dataset 1 (Gutenberg-8k)

Table 4 shows performance on Gutenberg-8k ($\sigma = 10$, min. phrase length 5). SPMF-based baselines (GST/SA+LCP, BIDE+) run out of memory beyond $N = 500$. Among stringology-native algorithms, two regimes are visible: at mid-range ($N \leq 1,000$), BloomSpan is $\approx 20\%$ faster than FHK/DFI as the CBF filter amortises; at $N \geq 3,000$, FHK/DFI become 10–20% faster because the surviving MCFP set grows super-linearly under the permissive threshold ($\approx 218k$ phrases at $N = 5,000$), making BloomSpan’s output-sensitive maximality phase the bottleneck. BloomSpan runs out of memory at $N = 8,000$ under this operating point, while FHK and DFI complete. The MCFP sets of BloomSpan, FHK, and DFI are sorted and compared at every $N \in \{100, 250, 500, 1,000, 1,500, 3,000, 5,000\}$; all sets are identical across all measured configurations, confirming no false negatives or false positives relative to two independently implemented exact baselines. We do not claim equivalence beyond $N = 5,000$ where no exact baseline completes simultaneously.

4.7 Results: Realistic Dataset 2 (Gutenberg-BookCorpus-Cleaned)

Table 5 and Figure 4 report total wall-clock time and peak heap on nine subsets (1,000–50,000 documents, $\approx 64\text{M}$ – 3.07B tokens). Each configuration used six

Table 5: BloomSpan vs. FHK vs. DFI on *Gutenberg-BookCorpus-Cleaned* ($\sigma = 20$, min. phrase length 10). Wall-clock time in seconds; peak heap memory in MB. Medians over five measured runs. Speedup = FHK / BloomSpan time. OOM = Out of Memory (40 GB heap).

N	BloomSpan		FHK		DFI		Speedup (FHK/BS)	MCFPs
	Time (s)	Mem (MB)	Time (s)	Mem (MB)	Time (s)	Mem (MB)		
1,000	1.73	5,095	19.08	6,591	20.17	9,138	11.0×	31
3,000	4.16	7,836	60.34	7,856	67.25	29,393	14.5×	120
5,000	6.64	6,977	101.33	11,238	OOM	OOM	15.3×	217
8,000	10.41	8,853	160.37	16,708	OOM	OOM	15.4×	375
10,000	13.59	12,044	199.34	20,764	OOM	OOM	14.7×	516
15,000	20.80	14,979	308.67	31,549	OOM	OOM	14.8×	1,034
20,000	27.67	20,839	–	–	OOM	OOM	–	1,964
30,000	44.61	29,385	662.24	38,461	OOM	OOM	14.8×	3,673
50,000	383.27	38,873	OOM	OOM	OOM	OOM	–	–

runs (one warm-up discarded); medians with IQR are reported. BloomSpan is 11–15× faster than FHK across every point where both complete (Table 5). The memory advantage grows with N : at $N = 15,000$, FHK uses 31.5 GB versus BloomSpan’s 15.0 GB. DFI’s $O(n \log n)$ RMQ sparse table consumes ≈ 29.4 GB already at $N = 3,000$ and fails with OOM at $N = 5,000$. FHK completes through $N = 30,000$ before exceeding the heap. BloomSpan completes all nine subsets including $N = 50,000$.

BloomSpan’s per-token mining cost on this operating point is approximately constant up to $N = 30,000$: measured in $\mu\text{s}/\text{token}$ for the table’s mining-only values, it sits in the range 0.021–0.027 across that range (a $\approx 28\times$ increase in input size), corresponding to near-linear scaling with a log-log slope of 0.98. At $N = 50,000$ the per-token cost jumps to $0.125 \mu\text{s}/\text{token}$ as the surviving pattern set reaches 12,244 MCFPs and the output-sensitive maximality-check phase becomes visible in the mining time; log-log slope over all nine points is 1.24. This super-linear point is nevertheless the only configuration that completes on a 3 B-token corpus under the 40 GB heap among all tested algorithms.

4.8 Topological Analysis of Datasets

The performance differences across datasets stem from measurable structural properties of each corpus. Let c denote the number of Lempel-Ziv phrases and n the total token count.

Synthetic Dataset 1 (Sparse Injections). Background words are entirely unique, yielding a nearly flat frequency distribution and a maximized phrase generation ratio ($c/n \approx 0.99$, $C_{LZ} \approx 24.4$ bits/token). Under these conditions the CBF operates at peak efficiency, pruning non-repeating background noise with zero false negatives and minimal memory overhead. BloomSpan therefore scales near-linearly and decisively outruns the gap-aware miners (BIDE+,




Fig. 4: Wall-clock time (left) and peak heap memory (right) on **Gutenberg-BookCorpus-Cleaned** ($\sigma = 20$, min. phrase length 10). DFI exhausts the heap beyond $N = 3,000$, *tokens* = 190M; FHK completes through $N = 30,000$, *tokens* = 1,7B.

GST/SA+LCP). Among stringology-native baselines, FHK is consistently faster: its LCP-interval traversal directly enumerates maximal substrings without a separate seed-filtering pass, so BloomSpan’s CBF lookup is pure overhead here. Before the crossover point, $N = 12,000$, BloomSpan’s advantage on this topology is memory. Because the background consists of entirely unique tokens, the phrase generation ratio is maximized ($c/n \approx 0.99$). In such a "flat" distribution, BloomSpan’s Counting Bloom Filter (CBF) provides less pruning benefit relative to its overhead, whereas FHK’s suffix-array structures handle the lack of deep branching more natively.

Synthetic Dataset 2 (Dense Nested Overlaps). The 300 heavily overlapping injections shift the topological advantage to suffix-based structures: FHK compresses identical overlapping prefixes and suffixes into a single LCP-interval traversal, neutralising the combinatorial explosion. BloomSpan’s CBF-based pruning loses much of its advantage because nearly every n -gram in the injected pattern family is frequent; the maximality-verification phase does most of the work. Despite this, BloomSpan finishes ahead of BIDE+ (Contiguous Maximal Mode) because BIDE+’s per-pattern redundancy tracking scales poorly with deeply nested overlaps.

Realistic Datasets (Project Gutenberg). Natural language provides a balanced, organic topology with a low phrase generation ratio ($c/n \approx 0.27$) and an LZ entropy rate of roughly 7.94 bits/token. This is the domain BloomSpan is designed to exploit. The Zipfian frequency distribution [24] typical of natural language produces a long tail of infrequent n -grams; the CBF prunes this tail effectively, eliminating 98% of candidate seeds at $\sigma = 20$ (Section 4.7). The moderate compressibility also indicates that patterns are not pathologically nested, preventing the exponential memory explosion that cripples GST/SA+LCP and BIDE+ on real text and allowing BloomSpan’s occupancy matrix to efficiently suppress redundancies. On Realistic Dataset 2 under a stricter operating point the output-sensitive phase stays sub-critical, and BloomSpan achieves the near-linear per-token scaling reported in Section 4.7.

5 Conclusion and Future Work

We presented BloomSpan, an algorithm for MCFP mining that combines a Counting Bloom Filter with prioritized depth-first expansion and a global occupancy bitmask, targeting natural-language corpora at strict support thresholds.

Evaluation on four datasets reveals clear algorithm-topology trade-offs. On both synthetic datasets (sparse injections and dense nested overlaps) FHK is the fastest in wall-clock time, with BloomSpan second; on sparse injections at large N BloomSpan uses the least memory among stringology-native miners. On Gutenberg-8k at the permissive threshold $\sigma = 10$, FHK/DFI are marginally faster at scale as BloomSpan’s data-dependent expansion phase becomes the bottleneck. At a strict-threshold large-scale configuration (Realistic Dataset 2, $\sigma = 20$, minimum length 10, up to 3 B tokens), BloomSpan is 11–15 $\times$ faster than FHK, scales near-linearly per token up to $N = 30,000$ (log-log slope 0.98), and is the only evaluated algorithm completing at $N = 50,000$ within a 40 GB heap. Classical gap-aware miners (BIDE+, VMSP, GST/SA+LCP) OOM at all but the smallest scales on realistic corpora.

Future directions include near-duplicate pattern support (edit-distance tolerance), a distributed implementation with a shared CBF, and validation on non-literary domains such as genomic sequences and system event logs.

References

1. Abouelhoda, M.I., Kurtz, S., Ohlebusch, E.: Replacing suffix trees with enhanced suffix arrays. *Journal of Discrete Algorithms* **2**(1), 53–86 (2004)
2. Agrawal, R., Srikant, R.: Mining sequential patterns. In: *Proceedings of the Eleventh International Conference on Data Engineering*. pp. 3–14. IEEE (1995)
3. Cormode, G., Muthukrishnan, S.: An improved data stream summary: the count-min sketch and its applications. *Journal of Algorithms* **55**(1), 58–75 (2005)
4. Dhaliwal, J., Puglisi, S.J., Turpin, A.: Practical efficient string mining. *IEEE Transactions on Knowledge and Data Engineering* **24**(4), 735–744 (2012). <https://doi.org/10.1109/TKDE.2010.242>
5. Fan, L., Cao, P., Almeida, J., Broder, A.Z.: Summary cache: a scalable wide-area web cache sharing protocol. *IEEE/ACM Transactions on Networking* **8**(3), 281–293 (2000)
6. Fhrozen: gutenber8k: A snapshot of 8544 english-language books from project gutenber. *Hugging Face Datasets* (2023), <https://huggingface.co/datasets/Fhrozen/gutenber8k>, accessed 2026
7. Fischer, J., Heun, V.: Fast frequent string mining using suffix arrays. In: *Annual Symposium on Combinatorial Pattern Matching*. pp. 155–166. Springer (2006)
8. Fischer, J., Heun, V., Kramer, S.: Fast frequent string mining using suffix arrays. In: *Proc. 5th IEEE International Conference on Data Mining (ICDM)*. pp. 609–612 (2005). <https://doi.org/10.1109/ICDM.2005.62>
9. Fischer, J., Heun, V., Kramer, S.: Fast frequent string mining using suffix arrays. *Data Mining and Knowledge Discovery* **18**(3), 439–479 (2009). <https://doi.org/10.1007/s10618-008-0122-1>

10. Fischer, J., Mäkinen, V., Välimäki, N.: Space-efficient string mining under frequency constraints. In: Proc. 8th IEEE International Conference on Data Mining (ICDM). pp. 193–202 (2008). <https://doi.org/10.1109/ICDM.2008.32>
11. Fournier-Viger, P., Lin, J.C.W., Kiran, R.U., Koh, Y.S., Thomas, R.: The SPMF open-source data mining library version 2. In: Joint European Conference on Machine Learning and Knowledge Discovery in Databases. pp. 36–40. Springer (2016)
12. Fournier-Viger, P., Wu, C.W., Gomariz, A., Tseng, V.S.: Vmsp: Efficient vertical mining of maximal sequential patterns. In: Canadian Conference on Artificial Intelligence. pp. 83–94. Springer (2014)
13. Guo, P., Umboh, S.W., Wirth, A., Zobel, J.: Online computation of string net frequency. In: String Processing and Information Retrieval — 31st International Symposium, SPIRE 2024. Lecture Notes in Computer Science, vol. 14899, pp. 159–173. Springer (2024). https://doi.org/10.1007/978-3-031-72200-4_12
14. incredible45: Gutenberg-bookcorpus-cleaned-data-english: a cleaned snapshot of the project gutenberg english book corpus. Hugging Face Datasets (2024), <https://huggingface.co/datasets/incredible45/Gutenberg-BookCorpus-Cleaned-Data-English>, accessed 2026
15. Kasai, T., Lee, G., Arimura, H., Arikawa, S., Park, K.: Linear-time longest-common-prefix computation in suffix arrays and its applications. In: Proc. 12th Annual Symposium on Combinatorial Pattern Matching (CPM). Lecture Notes in Computer Science, vol. 2089, pp. 181–192. Springer (2001). https://doi.org/10.1007/3-540-48194-X_17
16. Kügel, A., Ohlebusch, E.: A space efficient solution to the frequent string mining problem for many databases. Data Mining and Knowledge Discovery **17**(1), 24–38 (2008). <https://doi.org/10.1007/s10618-008-0110-5>
17. Manber, U., Myers, E.W.: Suffix arrays: A new method for on-line string searches. SIAM Journal on Computing **22**(5), 935–948 (1993). <https://doi.org/10.1137/0222058>
18. Navarro, G., Mäkinen, V.: Compressed full-text indexes. ACM Computing Surveys **39**(1), 2:1–2:61 (2007). <https://doi.org/10.1145/1216370.1216372>
19. Nong, G., Zhang, S., Chan, W.H.: Two efficient algorithms for linear time suffix array construction. IEEE Transactions on Computers **60**(10), 1471–1484 (2011). <https://doi.org/10.1109/TC.2010.188>
20. Ohlebusch, E., Büchler, T., Olbrich, J.: Faster computation of chinese frequent strings and their net frequencies. In: String Processing and Information Retrieval — 31st International Symposium, SPIRE 2024. Lecture Notes in Computer Science, vol. 14899, pp. 249–256. Springer (2024). https://doi.org/10.1007/978-3-031-72200-4_19
21. Pei, J., Han, J., Mortazavi-Asl, B., Pinto, H., Chen, Q., Dayal, U., Hsu, M.C.: Prefixspan: Mining sequential patterns efficiently by prefix-projected pattern growth. In: Proceedings of the 17th International Conference on Data Engineering. pp. 215–224. IEEE (2001)
22. Wang, J., Han, J.: Bide: Efficient mining of frequent closed sequences. In: Proceedings of the 20th International Conference on Data Engineering. pp. 79–90. IEEE (2004)
23. Weese, D., Schulz, M.H.: Efficient string mining under constraints via the deferred frequency index. In: Proc. 8th Industrial Conference on Advances in Data Mining (ICDM/ADMA). Lecture Notes in Computer Science, vol. 5077, pp. 374–388. Springer (2008). https://doi.org/10.1007/978-3-540-70720-2_29
24. Zipf, G.K.: Human Behavior and the Principle of Least Effort. Addison-Wesley, Cambridge, MA (1949)